

Improving convergence of stochastic gradient descent

Neural network – large general parametrization

How to train them: choose these parameters to represent given data?

By some gradient descent ... how to choose step size, pass saddles?

$$\theta^{t+1} = \theta^t - \eta g^t \quad \text{e.g. for} \quad g^t = \nabla_{\theta} F(\theta^t) \quad F(\theta) = \frac{1}{n} \sum_{i=1}^n L(x_i, \theta)$$

Gradient for entire (huge n) dataset, or online updates: mini-batches

Small batch – cheap, but large noise, how to extract statistical trends?

Linear or model parabola - modelling distance to extremum? (step size)

In one or multiple directions? especially near saddles

$\exp(\text{dim})$ saddles >> # minima at least from parameter permutations

There is some idealized

hidden probability distribution $p(x)$

e.g. of pictures of digits among bitmaps of given resolution, we want to label them

We need **parameters $\theta \in \mathbb{R}^D$**

minimizing loss/objective function:

$$\tilde{F}(\theta) = \int L(x, \theta) p(x) dx$$

But we know only samples $\{x_i\}_{i=1..n}$

Instead find local minimum of:

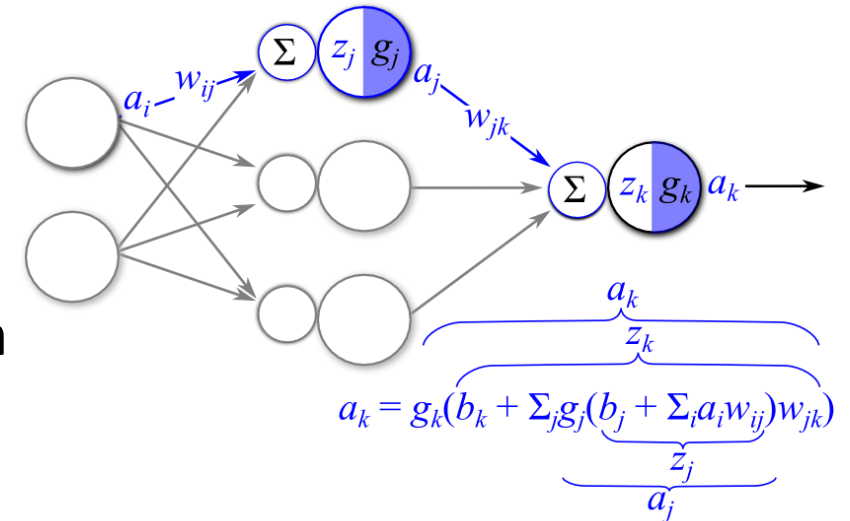
$$F(\theta) = \frac{1}{n} \sum_{i=1}^n L(x_i, \theta)$$

using **gradient sequence $\{\nabla_{\theta} F(\theta^t)\}_t$**


e.g. from backpropagation of errors.

Generalization: to prevent **overfitting**,
train on one subset,
test on separate validation set.


I. Forward-propagate Input Signal



II. Back-propagate Error Signals



III. Calculate Parameter Gradients



IV. Update Parameters

$$w_{ij} = w_{ij} - \eta (\partial E / \partial w_{ij})$$

$$w_{jk} = w_{jk} - \eta (\partial E / \partial w_{jk})$$

for learning rate η

We would like to minimize $F(\theta) = \frac{1}{n} \sum_{i=1}^n L(x_i, \theta)$ based on gradients:

- 1) **Batch/vanilla gradient descent**: using $g^t = \frac{1}{n} \sum_{i=1}^n \nabla_{\theta} L(x_i, \theta^t)$
 - 2) **Stochastic gradient descent (SGD)**: “online” $g^t = \nabla_{\theta} L(x_{it}, \theta^t)$
 - 3) **Mini-batch (~100) gradient descent**: using gradient from subsets
- Averaging over entire dataset (batch) – accurate but extremely costly,
over a subset – cheaper but noisy – average over time to real gradient
Let’s combine 2) and 3): gradients from size 1+ subsets, averaging to ~real


Symmetry of parameters – lower bound for number of local minima

It seems most of local minima have close value to global minimum here

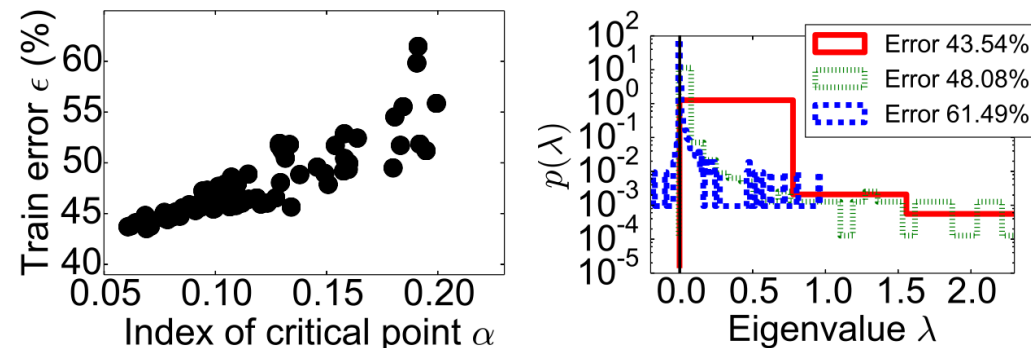
Usually much more saddles: $\binom{D}{\alpha D} \approx 2^{D h(\alpha)}$ assuming some randomness

α – percent of positive Hessian eigenvalues vs loss function:

MNIST





CIFAR-10




SGD ([overview](#)) – “**update** parameters **during** calculating gradient”

- Challenges from standard gradient descent:
oscillations in high curvature directions,
better [conjugate gradients](#) using also low curvature,
[saddles](#), also degenerated: with $\lambda_i = 0$
often large **plateaus** especially near saddles
- Choosing **step size**, their **schedule**?
- plus **huge dimension** and **noisy** gradients:
need to **extract statistical trends** ...



$(x, -y)$



SGD optimization: $g^t = \nabla_{\theta} F^t(\theta^t)$ noisy, averages to $\nabla_{\theta} F(\theta^t)$

Momentum - use **exponential moving average** of stochastic gradients:

$$v^t = \gamma v^{t-1} + (1 - \gamma) g^t = (1 - \gamma) g^t + \gamma(1 - \gamma) g^{t-1} + \dots$$


$$\theta^{t+1} = \theta^t - \eta v^t \quad \gamma \approx 0.9$$

Nesterov accelerated gradient (NAG):

“implicit Euler momentum”

$$v^t = \gamma v^{t-1} + \eta \nabla_{\theta} F^t(\theta^t - \gamma v^{t-1})$$

$$\theta^{t+1} = \theta^t - v^t$$



There are better ODE methods:


Higher order error = $O(h^{p+1})$

Like Runge-Kutta?

Should we go this way for SGD???

plot

One Step: Implicit vs. Explicit



Adagrad – larger updates for rare parameters (i), smaller for common

$$\theta_i^{t+1} = \theta_i^t - \frac{\eta}{\sqrt{G_i^t + \epsilon}} g_i^t \quad G_i^t = G_i^{t-1} + (g_i^t)^2 \quad \epsilon \approx 10^{-8}$$

RMSprop – Adagrad with exponential moving average instead of sum

$$\theta_i^{t+1} = \theta_i^t - \frac{\eta}{\sqrt{G_i^t + \epsilon}} g_i^t \quad G_i^t = \gamma G_i^{t-1} + (1 - \gamma)(g_i^t)^2$$

Adadelta – analogously estimate step size (diagonal Hessian approx.?)

$$\theta_i^{t+1} = \theta_i^t - \frac{\sqrt{\Delta_i^t + \epsilon}}{\sqrt{G_i^t + \epsilon}} g_i^t \quad \Delta_i^t = \gamma \Delta_i^{t-1} + (1 - \gamma)(\theta_i^t - \theta_i^{t-1})^2$$

Adam (18k citations): Adadelta + bias while starting exp. moving avg.

$$m_i^t = \beta_1 m_i^{t-1} + (1 - \beta_1) g_i^t \quad v^t = \beta_2 v^{t-1} + (1 - \beta_2)(g_i^t)^2$$
$$\theta^{t+1} = \theta^t - \frac{\eta}{\sqrt{v^t / (1 - (\beta_2)^t) + \epsilon}} \frac{m^t}{1 - (\beta_1)^t} \quad \beta_1 = 0.9, \beta_2 = 0.999$$

$$m_i^t = \beta_1 m_i^{t-1} + (1 - \beta_1) g_i^t \quad v^t = \beta_2 v^{t-1} + (1 - \beta_2) (g_i^t)^2$$

Adam:
$$\theta^{t+1} = \theta^t - \frac{\eta}{\sqrt{v^t / (1 - (\beta_2)^t)} + \epsilon} \frac{m^t}{1 - (\beta_1)^t}$$

AdaMax – Adam with maximum norm for stability (?)

$$\theta^{t+1} = \theta^t - \frac{\eta}{u_i^t} \frac{m^t}{1 - (\beta_1)^t} \quad u_i^t = \max(\beta_2 u_i^{t-1}, |g_i^t|)$$

Nadam – Adam + Nesterov (m^t estimated one step forward)


$$\theta^{t+1} = \theta^t - \frac{\eta}{\sqrt{v^t (1 - (\beta_2)^t)} + \epsilon} \left(\frac{\beta_1 m^t + (1 - \beta_1) g_t}{1 - (\beta_1)^t} \right)$$

AMSGrad – Adam often suboptimal

AdaBound

$$\hat{v}^t = \max(\hat{v}^{t-1}, v^t)$$

$$\theta^{t+1} = \theta^t - \frac{\eta}{\sqrt{\hat{v}^t} + \epsilon} m^t$$



Lots of heuristics, based on experiments, tasks of various specifics ...

They do not **estimate distance to extremum**, **trace only single direction**

Exponential number of saddles due to parameter permutation invariance

To quickly pass problematic **saddle** we could **model two parabolas** ...

Maybe let's try to go to **second order methods** ... not successful so far (?)

Newton-Raphson, $H > 0$: $\nabla_{\theta} F(\theta) \approx \nabla_{\theta} F(\theta^t) + H(\theta^t) \cdot (\theta - \theta^t)$

$\nabla_{\theta} F(\theta) \approx 0$ for “**natural gradient**”: $\theta - \theta^t = -H^{-1}(\theta^t) \cdot \nabla_{\theta} F(\theta^t)$

However, **huge dimensions** ($10^6 \dots 10^9$) and **noisy gradients** – need to **extract statistics**, **restrict Hessian**, **avoid inversion**, **avoid saddles** ...




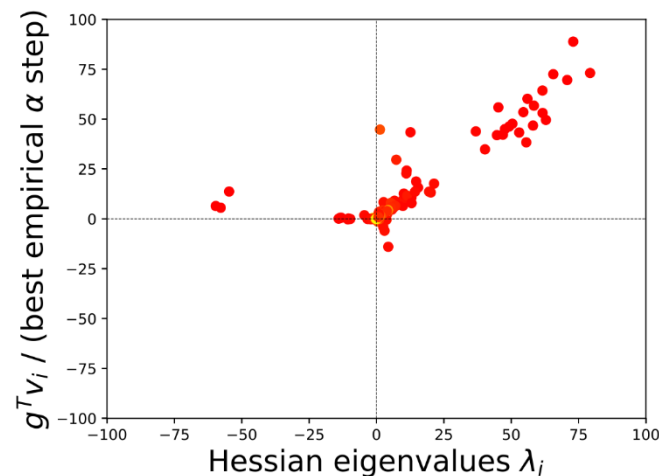

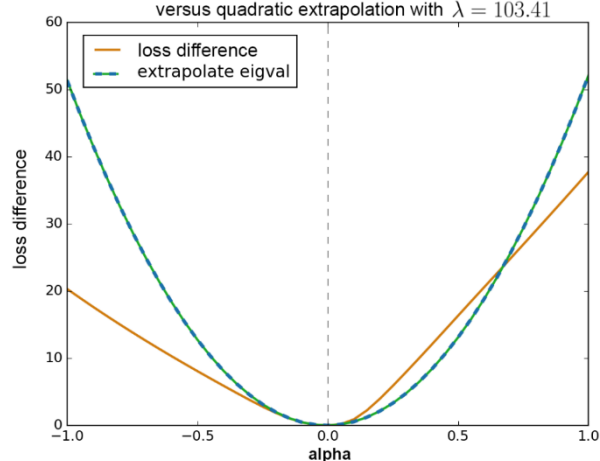

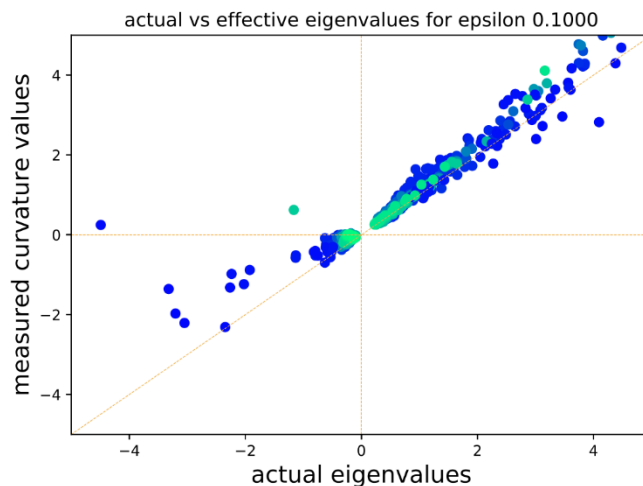

Let's start with simpler question (~“line search”):

How to handle 1D minimization asking for noisy derivatives?

$f = \lambda(x - p)^2/2$ $f' = \lambda(x - p)$ (expon. weight) **linear regression?**

Then: do it simultaneously for **a few directions**? How to explore them?

[arXiv:1902.02366](https://arxiv.org/abs/1902.02366) RMSprop on MNIST, $d = 3.3 \cdot 10^6$, $F(\theta + \alpha(g \cdot v_i)v_i)$, early: blue/red



Saddle-free Newton (~600cit.) [GitHub](#)

<https://arxiv.org/pdf/1406.2572>

$-|H|^{-1}g$ direction

$\lambda_i \rightarrow |\lambda_i|$

Lower dim. [Krylov](#): $\text{span}\{v, Hv, \dots, H^{k-1}v\}$

MSGD: minibatch SGD

Damped Newton: $-(H + \epsilon I)^{-1}g$

Algorithm 1 Approximate saddle-free Newton

Require: Function $f(\theta)$ to minimize

for $i = 1 \rightarrow M$ **do**

$\mathbf{V} \leftarrow k$ Lanczos vectors of $\frac{\partial^2 f}{\partial \theta^2}$

$\hat{f}(\alpha) \leftarrow g(\theta + \mathbf{V}\alpha)$

$|\hat{\mathbf{H}}| \leftarrow \left| \frac{\partial^2 \hat{f}}{\partial \alpha^2} \right|$ by using an eigen decomposition of $\hat{\mathbf{H}}$

for $j = 1 \rightarrow m$ **do**

$\mathbf{g} \leftarrow -\frac{\partial \hat{f}}{\partial \alpha}$

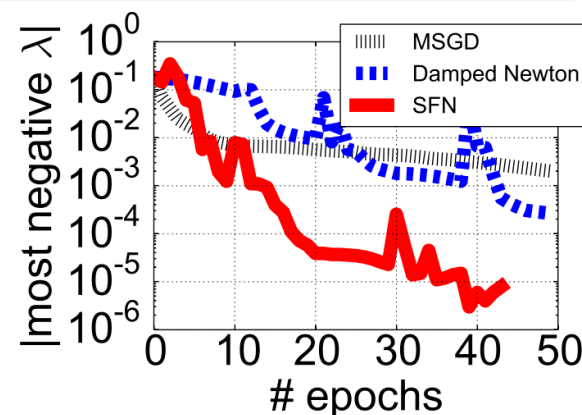


$\lambda \leftarrow \arg \min_{\lambda} \hat{f}(\mathbf{g}(|\hat{\mathbf{H}}| + \lambda \mathbf{I})^{-1})$

$\theta \leftarrow \theta + \mathbf{g}(|\hat{\mathbf{H}}| + \lambda \mathbf{I})^{-1} \mathbf{V}$




end for

end for

MNIST



CIFAR-10



Natural gradient attracts to saddle, there are usually $\exp(\text{dim})$ of them ...

Gauss-Newton method: assume $F(\boldsymbol{\theta}) = \sum_{i=1}^n (f_i(\boldsymbol{\theta}))^2$ “sample errors”

$$H_{jk} = \frac{\partial F}{\partial \theta_j \partial \theta_k} = 2 \sum_{i=1}^n \left(\frac{\partial f_i}{\partial \theta_j} \frac{\partial f_i}{\partial \theta_k} + f_i \frac{\partial^2 f_i}{\partial \theta_j \partial \theta_k} \right)$$

Assuming $\frac{\partial^2 f_i}{\partial \theta_j \partial \theta_k} = 0$, locally approximating f with linear functions:

$$H_{jk} \approx 2 \sum_{i=1}^n \frac{\partial f_i}{\partial \theta_j} \frac{\partial f_i}{\partial \theta_k} \quad \text{positive definite, only gradients needed (saddles?)}$$

$$\boldsymbol{\theta}^{t+1} = \boldsymbol{\theta}^t - (J_f^T J_f)^{-1} J_f \mathbf{f} \quad \left(= (J_f)^{-1} \mathbf{f} \text{ for } n = D \right) \quad J_f = \frac{\delta f_i}{\delta \theta_j}(\boldsymbol{\theta}^t)$$

Levenberg-Marquardt ($\lambda > 0$): $\boldsymbol{\theta}^{t+1} = \boldsymbol{\theta}^t - (J_f^T J_f + \lambda I)^{-1} J_f \mathbf{f}$

Other standard ways to **models Hessian**: **finite differences** (of gradients),
or **backpropagating** – often dropping 2nd derivative of activation function

Classical second order: conjugated gradients

“unbend to Hessian eigenbasis” $\langle u, v \rangle_H := u^T H v$

(v_1, \dots, v_d) : $\langle v_i, v_j \rangle_H = 0$ for $i \neq j$

$$v_k = r_k - \sum_{i < k} \frac{v_i^T H r_k}{v_i^T H v_i} v_i \quad \alpha_k = \frac{v_k^T r_k}{v_k^T H v_k}$$

$$v_{k+1} = x_k + \alpha_k v_k$$

Learning: truncated Newton or

nonlinear conjugated gradients (NCG):

Line search $\operatorname{argmin}_{\alpha > 0} F(x_t - \alpha v_t)$ toward


$$v_t = \beta_t v_{t-1} - \nabla F(\theta^t) \quad (+ \text{ reset sometimes})$$

satisfying $v_t^T H v_{t-1} \approx 0$. From LeCun et al. “Efficient BackProp” (1998):

Fletcher and Reeves (1964), or Polak and Ribierre (1969):

$$\beta_t = \frac{\nabla F(\theta^t)^T \nabla F(\theta^t)}{\nabla F(\theta^{t-1})^T \nabla F(\theta^{t-1})}$$

$$\beta_t = \frac{(\nabla F(\theta^t) - \nabla F(\theta^{t-1}))^T \nabla F(\theta^t)}{\nabla F(\theta^{t-1})^T \nabla F(\theta^{t-1})}$$



Quasi-Newton – instead of inverting Hessian, update approximation of H^{-1}

BFGS (Broyden-Fletcher-Goldfarb-Shanno ~1970), positive definite H

Secant condition that gradients agree: $H_n^{-1} \Delta g_n = \Delta x_n$

$$\nabla F_n(x_n) = g_n, \quad \nabla F_n(x_{n-1}) = g_{n-1} \quad \Rightarrow \quad H_n(x_n - x_{n-1}) = (g_n - g_{n-1})$$

Find $\operatorname{argmin}_{H^{-1}=(H^{-1})^T} \|H^{-1} - H_n^{-1}\|_F^2$ s.t. $H_n^{-1} \Delta g_n = \Delta x_n$ getting:

$$H_{n+1}^{-1} = \left(I - \frac{\Delta g_n (\Delta x_n)^T}{\Delta g_n \cdot \Delta x_n} \right) H_n^{-1} \left(I - \frac{\Delta x_n (\Delta g_n)^T}{\Delta g_n \cdot \Delta x_n} \right) + \frac{\Delta x_n (\Delta x_n)^T}{\Delta g_n \cdot \Delta x_n}$$

$$x_{n+1} = x_n - H_{n+1}^{-1} g_n$$

or **line search**: $\operatorname{argmin}_{\alpha > 0} F(x_n - \alpha H_{n+1}^{-1} g_n)$

L-BFGS: limited memory: store $m \approx 10$ last $\Delta x, \Delta g$ instead of huge H^{-1}

“two loop recursion”: $i = n \dots \searrow \dots n - m \dots \nearrow \dots n$ using $H_{n-m}^{-1} \approx I$

Rough approximation, numerical problem with **noisy gradients** in SGD

K-FAC (2015): “Kronecker-factored approximate curvature”, [slides](#)

Natural gradient approximating full Hessian with **block-diagonal** full Hessian inside **layers**, no inter-layer correlations (tri-block-diagonal...)

Huge cost: $\sim \text{dimension}^2 / \# \text{layers}$ noisy estimations ... **saddles** ($\approx H > 0$)

Maximizing likelihood with density f_θ : $F(\theta) = \frac{1}{n} \sum_{i=1}^n \ln(f_\theta(x_i))$

$$\sum_{ij} u_i E \left[\left(\frac{\partial \ln f_\theta(X)}{\partial \theta_i} \frac{\partial \ln f_\theta(X)}{\partial \theta_j} \right) \right] u_j = E \left[\left(\sum_i u_i \frac{\partial \ln f_\theta(X)}{\partial \theta_i} \right)^2 \right] > 0$$


Fisher information

is positive definite

describes parameter
dependence/certainty


\sim Gauss-Newton

Hessian approx.



H_k

\approx



\hat{H}_k

A fast natural Newton method (2010) – TONGA introduction + repair

Use (~PCA) correlation of recent gradients instead of Hessian

Assume g : “true” gradient, \hat{g} : “empirical gradient” from Gaussian:

$$\hat{g}|g \sim \mathcal{N}(g, C/n) \quad \text{for centered covariance matrix } C$$

$$C = \int_x \left(\frac{\partial f(\theta, x)}{\partial \theta} - g \right) \left(\frac{\partial f(\theta, x)}{\partial \theta} - g \right)^T p(x) dx \quad \text{isotropic } g \sim \mathcal{N}(0, \sigma^2 I)$$

$$\hat{g}|g \sim \mathcal{N} \left(\left[I + \frac{C}{n\sigma^2} \right]^{-1} \hat{g}, [nC^{-1} + \sigma^{-2}I]^{-1} \right) \quad \text{saddle?}$$

$$\Delta\theta \propto - \left[I + \frac{\hat{C}}{n\sigma^2} \right]^{-1} \hat{g} \quad \text{for } C \approx \hat{C} = \frac{1}{n} \sum_{i=1}^n \left(\frac{\partial f(\theta, x_i)}{\partial \theta} - g \right) \left(\frac{\partial f(\theta, x_i)}{\partial \theta} - g \right)^T$$

2010 improvement (full dimension?):

$$\text{Assuming } F(\theta) \approx \frac{1}{2} (\theta - \theta^*)^T H (\theta - \theta^*) \quad \text{prior } g \sim \mathcal{N}(0, \sigma^2 H)$$

$$\Delta\theta \propto - \left[I + \frac{H^{-1} \hat{C} H^{-1}}{n\sigma^2} \right]^{-1} \hat{g} \quad \text{Hessian } H \text{ from Quasi-Newton}$$

Uncentered covariance matrix from exponential moving average ...

TONGA (2008): reduced dimension, Le Roux, Bengio, Manzagol

non-centered covariance matrix from exponential moving average:

$$C_t = \gamma \hat{C}_t + g_t g_t^T \approx X_t X_t^T$$

X_t is $D \times d$ for some $d \ll D$

(standard (centered) covariance matrix: $\frac{1}{n} \sum_{i=1}^n (g_i - \bar{g})(g_i - \bar{g})^T$)

Regularized **natural gradient**:

$$v_t = (C_t + \lambda I)^{-1} g_t = X_t (X_t^T X_t + \lambda I)^{-1} y_t \quad \text{in } O(Dd + d^3) \text{ time}$$

\hat{C}_t – low rank approximation of C – keep only $k < d$ eigenvalues

$$G_t = X_t X_t^T = V D V^T$$

$$C_t = (X_t V D^{-1/2}) D (X_t V D^{-1/2})^T$$

eigendecomposition at cost $O(kd^2 + Ddk)$ every few steps

... maybe let's try to directly model and update only what we really need ...

Update local parametrization (saddles), put eigendecomposition in iteration

Online gradient linear regression: let's start with 1D fixed parabola:

$$f(\theta) = \frac{1}{2} \lambda (\theta - p)^2 \quad \text{from noisy} \quad g^t \approx f'(\theta^t) = \lambda(\theta^t - p)$$

Least-square linear regression: $\arg \min_{\lambda, p} \sum_t w^t (g^t - \lambda(\theta^t - p))^2$

$$\lambda = \frac{\overline{g\theta} - \bar{g} \cdot \bar{\theta}}{\overline{\theta^2} - \bar{\theta}^2} \quad p = \frac{\lambda \bar{\theta} - \bar{g}}{\lambda} \quad \text{e.g. } \overline{g\theta} = \frac{1}{T} \sum_t g^t \theta^t$$

For $\bar{g}, \bar{\theta^2}, \overline{g\theta}, \overline{\theta^2}$ $w^t = 1/T$ averages

In general (non-parabola) case: use **exponential moving average**

$$w^t \sim \beta^{-t} \text{ online averages, e.g. } \overline{g\theta}^t = \beta \overline{g\theta}^{t-1} + (1 - \beta) g^t \theta^t$$

In D dimensions: do it in a few $d \ll D$ statistically relevant $(v_i)_{i=1..d}$

Attract ($\lambda > 0$) **or repel** ($\lambda < 0$) correspondingly to handle **saddles**

$$\theta \leftarrow \theta + \alpha \sum_{i=1}^d \text{sign}(\lambda_i) (p_i - \theta \cdot v_i) v_i \quad \text{proper optimization step}$$

Maintain \approx diagonal Hessian by periodic diagonalization, orthogonalization

$$O^T \Lambda O = H = \left(\overline{g\theta} - \bar{g} \bar{\theta}^T \right) \left(\overline{\theta\theta} - \bar{\theta} \bar{\theta}^T \right)^{-1}, \text{ e.g. } \overline{\theta\theta}_{ij} = \overline{\theta_i \theta_j}$$

Rotate subspace (online) to explore recent relevant directions

Many tough questions

- Should we ask for **gradients**, or maybe (also?) **values**, **2nd derivatives**?
Gradients: nice compromise – suggests direction, cheaper than Hessian
- **How many directions** to model? One (now) ... a few ... all (full Hessian)?
- How to **choose interesting directions** based on recent gradients?
- Online: **mini-batch size**? Step cost vs uncertainty?
- Strengthening **rarely represented coordinates** like $g_i / \langle g_i^2 \rangle^{1/2}$?
- How to efficiently pass **plateaus, saddles**?
- How to efficiently handle noise – **extract statistical trends**?
- How frequent are saddles? (much more than minima)
- How bad are **positive Hessian approximations**? $F = \sum_i (f_i)^2$ only?
aren't they attracting to saddles?

Simpler warmup question: how to optimally handle 1D problem?

... how to optimize problem to reduce iteration number...? [To one?](#)