

# More on Convolutional Neural Networks

Romuald A. Janik

Data science/machine learning journal club at WFAIS UJ

## Outline

### **Convolutional Neural Networks (CNN)**

#### **Going deep**

#### **CNN example: MNIST**

#### **Image classification datasets**

#### **Recent CNN architectures**

- Inception/Xception

- Residual networks

- Densely connected convolutional networks

#### **Interlude — fooling deep CNN's**

#### **CNN for segmentation**

- Fully Convolutional Networks — U-Net

#### **Conclusions**

# Neural networks

## What are neural networks?

- ▶ Each neuron with inputs  $x_k$  from previous layer has the output  $y_i$  given by

$$y_i = f(W_{ik}x_k + b_i)$$


- ▶ The coefficients  $W_{ik}$  (*weights*) and  $b_i$  (*bias*) are determined during training
- ▶  $f(\cdot)$  is a **nonlinear** 'activation function'

$$f(x) = \tanh x \longrightarrow \underbrace{f(x) = \max(0, x)}_{\text{ReLU}}$$

- ▶ All neurons of a layer are connected to all neurons of the previous one
- ▶ This is good for general situations – but has lots of parameters (especially for deep networks)!
- ▶ Suboptimal for images

# Neural networks

## What are neural networks?



- ▶ Each neuron with inputs  $x_k$  from previous layer has the output  $y_i$  given by

$$y_i = f(W_{ik}x_k + b_i)$$


- ▶ The coefficients  $W_{ik}$  (weights) and  $b_i$  (bias) are determined during training
- ▶  $f(\cdot)$  is a **nonlinear** 'activation function'

$$f(x) = \tanh x \rightarrow \underbrace{f(x) = \max(0, x)}_{\text{ReLU}}$$

- ▶ All neurons of a layer are connected to all neurons of the previous one
- ▶ This is good for general situations – but has lots of parameters (especially for deep networks)!
- ▶ Suboptimal for images

# Neural networks

## What are neural networks?



- ▶ Each neuron with inputs  $x_k$  from previous layer has the output  $y_i$  given by

$$y_i = f(W_{ik}x_k + b_i)$$


- ▶ The coefficients  $W_{ik}$  (weights) and  $b_i$  (bias) are determined during training
- ▶  $f(\cdot)$  is a **nonlinear** 'activation function'

$$f(x) = \tanh x \rightarrow \underbrace{f(x) = \max(0, x)}_{\text{ReLU}}$$

- ▶ All neurons of a layer are connected to all neurons of the previous one
- ▶ This is good for general situations – but has lots of parameters (especially for deep networks)!
- ▶ Suboptimal for images

# Neural networks

## What are neural networks?



- ▶ Each neuron with inputs  $x_k$  from previous layer has the output  $y_i$  given by

$$y_i = f(W_{ik}x_k + b_i)$$


- ▶ The coefficients  $W_{ik}$  (weights) and  $b_i$  (bias) are determined during training
- ▶  $f(\cdot)$  is a **nonlinear** 'activation function'

$$f(x) = \tanh x \rightarrow \underbrace{f(x) = \max(0, x)}_{\text{ReLU}}$$

- ▶ All neurons of a layer are connected to all neurons of the previous one
- ▶ This is good for general situations – but has lots of parameters (especially for deep networks)!
- ▶ Suboptimal for images

# Neural networks

## What are neural networks?



- ▶ Each neuron with inputs  $x_k$  from previous layer has the output  $y_i$  given by

$$y_i = f(W_{ik}x_k + b_i)$$


- ▶ The coefficients  $W_{ik}$  (weights) and  $b_i$  (bias) are determined during training
- ▶  $f(\cdot)$  is a **nonlinear** 'activation function'

$$f(x) = \tanh x \rightarrow \underbrace{f(x) = \max(0, x)}_{\text{ReLU}}$$

- ▶ All neurons of a layer are connected to all neurons of the previous one
- ▶ This is good for general situations – but has lots of parameters (especially for deep networks)!
- ▶ Suboptimal for images

# Neural networks

## What are neural networks?



- ▶ Each neuron with inputs  $x_k$  from previous layer has the output  $y_i$  given by

$$y_i = f(W_{ik}x_k + b_i)$$

- ▶ The coefficients  $W_{ik}$  (weights) and  $b_i$  (bias) are determined during training
- ▶  $f(\cdot)$  is a **nonlinear** 'activation function'


$$f(x) = \tanh x \rightarrow \underbrace{f(x) = \max(0, x)}_{\text{ReLU}}$$

- ▶ All neurons of a layer are connected to all neurons of the previous one
- ▶ This is good for general situations – but has lots of parameters (especially for deep networks)!
- ▶ Suboptimal for images



# Neural networks

## What are neural networks?



- ▶ Each neuron with inputs  $x_k$  from previous layer has the output  $y_i$  given by

$$y_i = f(W_{ik}x_k + b_i)$$


- ▶ The coefficients  $W_{ik}$  (weights) and  $b_i$  (bias) are determined during training
- ▶  $f(\cdot)$  is a **nonlinear** 'activation function'

$$f(x) = \tanh x \rightarrow \underbrace{f(x) = \max(0, x)}_{\text{ReLU}}$$

- ▶ All neurons of a layer are connected to all neurons of the previous one
- ▶ This is good for general situations – but has lots of parameters (especially for deep networks)!
- ▶ Suboptimal for images

# Neural networks

## What are neural networks?



- ▶ Each neuron with inputs  $x_k$  from previous layer has the output  $y_i$  given by

$$y_i = f(W_{ik}x_k + b_i)$$


- ▶ The coefficients  $W_{ik}$  (weights) and  $b_i$  (bias) are determined during training
- ▶  $f(\cdot)$  is a **nonlinear** 'activation function'

$$f(x) = \tanh x \longrightarrow \underbrace{f(x) = \max(0, x)}_{\text{ReLU}}$$

- ▶ All neurons of a layer are connected to all neurons of the previous one
- ▶ This is good for general situations – but has lots of parameters (especially for deep networks)!
- ▶ Suboptimal for images

# Neural networks

## What are neural networks?



- ▶ Each neuron with inputs  $x_k$  from previous layer has the output  $y_i$  given by

$$y_i = f(W_{ik}x_k + b_i)$$


- ▶ The coefficients  $W_{ik}$  (weights) and  $b_i$  (bias) are determined during training
- ▶  $f(\cdot)$  is a **nonlinear** 'activation function'

$$f(x) = \tanh x \rightarrow \underbrace{f(x) = \max(0, x)}_{\text{ReLU}}$$

- ▶ All neurons of a layer are connected to all neurons of the previous one
- ▶ This is good for general situations – but has lots of parameters (especially for deep networks)!
- ▶ Suboptimal for images

# Neural networks

## What are neural networks?



- ▶ Each neuron with inputs  $x_k$  from previous layer has the output  $y_i$  given by

$$y_i = f(W_{ik}x_k + b_i)$$

- ▶ The coefficients  $W_{ik}$  (weights) and  $b_i$  (bias) are determined during training
- ▶  $f(\cdot)$  is a **nonlinear** 'activation function'

$$f(x) = \tanh x \longrightarrow \underbrace{f(x) = \max(0, x)}_{\text{ReLU}}$$

- ▶ All neurons of a layer are connected to all neurons of the previous one
- ▶ This is good for general situations – but has lots of parameters (especially for deep networks)!
- ▶ Suboptimal for images

# Convolutional Neural Networks (CNN)

## Main ideas:

1. Preserve the 2D structure of the input layer
2. Connect neurons in the next layer only to a block of say  $3 \times 3$  neurons in the previous layer
3. For all these connections take the same weights

from [http://deeplearning.stanford.edu/wiki/index.php/Feature\\_extraction\\_using\\_convolution](http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution)

# Convolutional Neural Networks (CNN)

## Main ideas:

1. Preserve the 2D structure of the input layer
2. Connect neurons in the next layer only to a block of say  $3 \times 3$  neurons in the previous layer
3. For all these connections take the same weights

from [http://deeplearning.stanford.edu/wiki/index.php/Feature\\_extraction\\_using\\_convolution](http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution)

# Convolutional Neural Networks (CNN)

## Main ideas:

1. Preserve the 2D structure of the input layer
2. Connect neurons in the next layer only to a block of say  $3 \times 3$  neurons in the previous layer
3. For all these connections take the same weights

from [http://deeplearning.stanford.edu/wiki/index.php/Feature\\_extraction\\_using\\_convolution](http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution)

# Convolutional Neural Networks (CNN)

## Main ideas:

1. Preserve the 2D structure of the input layer
2. Connect neurons in the next layer only to a block of say  $3 \times 3$  neurons in the previous layer
3. For all these connections take the same weights

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved  
Feature

from [http://deeplearning.stanford.edu/wiki/index.php/Feature\\_extraction\\_using\\_convolution](http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution)



# Convolutional Neural Networks (CNN)

## Main ideas:

1. Preserve the 2D structure of the input layer
2. Connect neurons in the next layer only to a block of say  $3 \times 3$  neurons in the previous layer
3. For all these connections take the same weights

1	1 <sub><math>\times 1</math></sub>	1 <sub><math>\times 0</math></sub>	0 <sub><math>\times 1</math></sub>	0
0	1 <sub><math>\times 0</math></sub>	1 <sub><math>\times 1</math></sub>	1 <sub><math>\times 0</math></sub>	0
0	0 <sub><math>\times 1</math></sub>	1 <sub><math>\times 0</math></sub>	1 <sub><math>\times 1</math></sub>	1
0	0	1	1	0
0	1	1	0	0

Image

4	3	

Convolved  
Feature

from [http://deeplearning.stanford.edu/wiki/index.php/Feature\\_extraction\\_using\\_convolution](http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution)

# Convolutional Neural Networks (CNN)

## Main ideas:

1. Preserve the 2D structure of the input layer
2. Connect neurons in the next layer only to a block of say  $3 \times 3$  neurons in the previous layer
3. For all these connections take the same weights

1	1	1 <sub><math>\times 1</math></sub>	0 <sub><math>\times 0</math></sub>	0 <sub><math>\times 1</math></sub>
0	1	1 <sub><math>\times 0</math></sub>	1 <sub><math>\times 1</math></sub>	0 <sub><math>\times 0</math></sub>
0	0	1 <sub><math>\times 1</math></sub>	1 <sub><math>\times 0</math></sub>	1 <sub><math>\times 1</math></sub>
0	0	1	1	0
0	1	1	0	0

Image

4	3	4

Convolved  
Feature

from [http://deeplearning.stanford.edu/wiki/index.php/Feature\\_extraction\\_using\\_convolution](http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution)

# Convolutional Neural Networks (CNN)

## Main ideas:

1. Preserve the 2D structure of the input layer
2. Connect neurons in the next layer only to a block of say  $3 \times 3$  neurons in the previous layer
3. For all these connections take the same weights

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4	3	4
2		

Convolved  
Feature

from [http://deeplearning.stanford.edu/wiki/index.php/Feature\\_extraction\\_using\\_convolution](http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution)

# Convolutional Neural Networks (CNN)

## Main ideas:

1. Preserve the 2D structure of the input layer
2. Connect neurons in the next layer only to a block of say  $3 \times 3$  neurons in the previous layer
3. For all these connections take the same weights

1	1	1	0	0
0	1 <sub><math>\times 1</math></sub>	1 <sub><math>\times 0</math></sub>	1 <sub><math>\times 1</math></sub>	0
0	0 <sub><math>\times 0</math></sub>	1 <sub><math>\times 1</math></sub>	1 <sub><math>\times 0</math></sub>	1
0	0 <sub><math>\times 1</math></sub>	1 <sub><math>\times 0</math></sub>	1 <sub><math>\times 1</math></sub>	0
0	1	1	0	0

Image

4	3	4
2	4	

Convolved  
Feature

from [http://deeplearning.stanford.edu/wiki/index.php/Feature\\_extraction\\_using\\_convolution](http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution)

# Convolutional Neural Networks (CNN)

## Main ideas:

1. Preserve the 2D structure of the input layer
2. Connect neurons in the next layer only to a block of say  $3 \times 3$  neurons in the previous layer
3. For all these connections take the same weights

1	1	1	0	0
0	1	1 <sub><math>\times 1</math></sub>	1 <sub><math>\times 0</math></sub>	0 <sub><math>\times 1</math></sub>
0	0	1 <sub><math>\times 0</math></sub>	1 <sub><math>\times 1</math></sub>	1 <sub><math>\times 0</math></sub>
0	0	1 <sub><math>\times 1</math></sub>	1 <sub><math>\times 0</math></sub>	0 <sub><math>\times 1</math></sub>
0	1	1	0	0

Image

4	3	4
2	4	3

Convolved  
Feature

from [http://deeplearning.stanford.edu/wiki/index.php/Feature\\_extraction\\_using\\_convolution](http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution)

# Convolutional Neural Networks (CNN)

## Main ideas:

1. Preserve the 2D structure of the input layer
2. Connect neurons in the next layer only to a block of say  $3 \times 3$  neurons in the previous layer
3. For all these connections take the same weights

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4	3	4
2	4	3
2		

Convolved  
Feature

from [http://deeplearning.stanford.edu/wiki/index.php/Feature\\_extraction\\_using\\_convolution](http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution)

# Convolutional Neural Networks (CNN)

## Main ideas:

1. Preserve the 2D structure of the input layer
2. Connect neurons in the next layer only to a block of say  $3 \times 3$  neurons in the previous layer
3. For all these connections take the same weights

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4	3	4
2	4	3
2	3	

Convolved  
Feature

from [http://deeplearning.stanford.edu/wiki/index.php/Feature\\_extraction\\_using\\_convolution](http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution)

# Convolutional Neural Networks (CNN)

## Main ideas:

1. Preserve the 2D structure of the input layer
2. Connect neurons in the next layer only to a block of say  $3 \times 3$  neurons in the previous layer
3. For all these connections take the same weights

1	1	1	0	0
0	1	1	1	0
0	0	1 <sub><math>\times 1</math></sub>	1 <sub><math>\times 0</math></sub>	1 <sub><math>\times 1</math></sub>
0	0	1 <sub><math>\times 0</math></sub>	1 <sub><math>\times 1</math></sub>	0 <sub><math>\times 0</math></sub>
0	1	1 <sub><math>\times 1</math></sub>	0 <sub><math>\times 0</math></sub>	0 <sub><math>\times 1</math></sub>

Image

4	3	4
2	4	3
2	3	4

Convolved  
Feature

from [http://deeplearning.stanford.edu/wiki/index.php/Feature\\_extraction\\_using\\_convolution](http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution)



# Convolutional Neural Networks (CNN)

- ▶ Motivation from visual system:
  1. First layers detect simple features – like edges at various angles (convolution weights are also called filters)
  2. This information is composed into higher level features in deeper layers of the network
- ▶ We will have multiple filters at each hidden layer

from <http://cs231n.github.io/convolutional-networks/>

# Convolutional Neural Networks (CNN)

- ▶ Motivation from visual system:
  1. First layers detect simple features – like edges at various angles (convolution weights are also called filters)
  2. This information is composed into higher level features in deeper layers of the network
- ▶ We will have multiple filters at each hidden layer

from <http://cs231n.github.io/convolutional-networks/>

# Convolutional Neural Networks (CNN)

- ▶ Motivation from visual system:
  1. First layers detect simple features – like edges at various angles (convolution weights are also called filters)
  2. This information is composed into higher level features in deeper layers of the network
- ▶ We will have multiple filters at each hidden layer

from <http://cs231n.github.io/convolutional-networks/>

# Convolutional Neural Networks (CNN)

- ▶ Motivation from visual system:
  1. First layers detect simple features – like edges at various angles (convolution weights are also called filters)
  2. This information is composed into higher level features in deeper layers of the network
- ▶ We will have multiple filters at each hidden layer

from <http://cs231n.github.io/convolutional-networks/>

# Convolutional Neural Networks (CNN)

- ▶ Motivation from visual system:
  1. First layers detect simple features – like edges at various angles (convolution weights are also called filters)
  2. This information is composed into higher level features in deeper layers of the network
- ▶ We will have multiple filters at each hidden layer

from <http://cs231n.github.io/convolutional-networks/>

# Convolutional Neural Networks (CNN)

- ▶ Motivation from visual system:
  1. First layers detect simple features – like edges at various angles (convolution weights are also called filters)
  2. This information is composed into higher level features in deeper layers of the network
- ▶ We will have multiple filters at each hidden layer

from <http://cs231n.github.io/convolutional-networks/>

# Convolutional Neural Networks (CNN)

- ▶ Motivation from visual system:
  1. First layers detect simple features – like edges at various angles (convolution weights are also called filters)
  2. This information is composed into higher level features in deeper layers of the network
- ▶ We will have multiple filters at each hidden layer



from <http://cs231n.github.io/convolutional-networks/>

# Convolutional Neural Networks (CNN)

**Sample structure of a CNN:**


Implement translational invariance

→ pooling layers



# Convolutional Neural Networks (CNN)

## Sample structure of a CNN:




Implement translational invariance

→ pooling layers

# Convolutional Neural Networks (CNN)

## Sample structure of a CNN:



## Implement translational invariance

→ **pooling layers**

# Convolutional Neural Networks (CNN)

Max-pooling layer


from <http://cs231n.github.io/convolutional-networks/>

Less popular: Average-pooling

**Pooling layers realize approximate translational invariance**

# Convolutional Neural Networks (CNN)

## Max-pooling layer




from <http://cs231n.github.io/convolutional-networks/>

Less popular: Average-pooling

Pooling layers realize approximate translational invariance

# Convolutional Neural Networks (CNN)

## Max-pooling layer



from <http://cs231n.github.io/convolutional-networks/>

Less popular: Average-pooling

**Pooling layers realize approximate translational invariance**

# Convolutional Neural Networks (CNN)

In a machine learning model we can distinguish:

- ▶ **Parameters** – learned during training (e.g. weights of the connections between neurons)
- ▶ **Hyperparameters** – fixed *a-priori* – define the structure of the network:

For a CNN:

1. number and type of layers
2. for each convolutional layer: filter size, number of channels/filters, stride (stride=2 in the figure below)
3. Additional layers: dropout/batch normalization/flatten/fully connected..

## Convolutional Neural Networks (CNN)

In a machine learning model we can distinguish:

- ▶ **Parameters** – learned during training (e.g. weights of the connections between neurons)
- ▶ **Hyperparameters** – fixed *a-priori* – define the structure of the network:

For a CNN:

1. number and type of layers
2. for each convolutional layer: filter size, number of channels/filters, stride (stride=2 in the figure below)
3. Additional layers: dropout/batch normalization/flatten/fully connected..

## Convolutional Neural Networks (CNN)

In a machine learning model we can distinguish:

- ▶ **Parameters** – learned during training (e.g. weights of the connections between neurons)
- ▶ **Hyperparameters** – fixed *a-priori* – define the structure of the network:

For a CNN:

1. number and type of layers
2. for each convolutional layer: filter size, number of channels/filters, stride (stride=2 in the figure below)
3. Additional layers: dropout/batch normalization/flatten/fully connected..



## Convolutional Neural Networks (CNN)

In a machine learning model we can distinguish:

- ▶ **Parameters** – learned during training (e.g. weights of the connections between neurons)
- ▶ **Hyperparameters** – fixed *a-priori* – define the structure of the network:

For a CNN:

1. number and type of layers
2. for each convolutional layer: filter size, number of channels/filters, stride (stride=2 in the figure below)
3. Additional layers: dropout/batch normalization/flatten/fully connected..

## Convolutional Neural Networks (CNN)

In a machine learning model we can distinguish:

- ▶ **Parameters** – learned during training (e.g. weights of the connections between neurons)
- ▶ **Hyperparameters** – fixed *a-priori* – define the structure of the network:

For a CNN:

1. number and type of layers
2. for each convolutional layer: filter size, number of channels/filters, stride (stride=2 in the figure below)
3. Additional layers: dropout/batch normalization/flatten/fully connected..

## Convolutional Neural Networks (CNN)

In a machine learning model we can distinguish:

- ▶ **Parameters** – learned during training (e.g. weights of the connections between neurons)
- ▶ **Hyperparameters** – fixed *a-priori* – define the structure of the network:

For a CNN:

1. number and type of layers
2. for each convolutional layer: filter size, number of channels/filters, stride (stride=2 in the figure below)
3. Additional layers: dropout/batch normalization/flatten/fully connected..


## Convolutional Neural Networks (CNN)

In a machine learning model we can distinguish:

- ▶ **Parameters** – learned during training (e.g. weights of the connections between neurons)
- ▶ **Hyperparameters** – fixed *a-priori* – define the structure of the network:

For a CNN:

1. number and type of layers
2. for each convolutional layer: filter size, number of channels/filters, stride (stride=2 in the figure below)
3. Additional layers: dropout/batch normalization/flatten/fully connected..



Image

Convolved  
Feature


## Convolutional Neural Networks (CNN)

In a machine learning model we can distinguish:

- ▶ **Parameters** – learned during training (e.g. weights of the connections between neurons)
- ▶ **Hyperparameters** – fixed *a-priori* – define the structure of the network:

For a CNN:

1. number and type of layers
2. for each convolutional layer: filter size, number of channels/filters, stride (stride=2 in the figure below)
3. Additional layers: dropout/batch normalization/flatten/fully connected..




## Convolutional Neural Networks (CNN)

In a machine learning model we can distinguish:

- ▶ **Parameters** – learned during training (e.g. weights of the connections between neurons)
- ▶ **Hyperparameters** – fixed *a-priori* – define the structure of the network:

For a CNN:

1. number and type of layers
2. for each convolutional layer: filter size, number of channels/filters, stride (stride=2 in the figure below)
3. Additional layers: dropout/batch normalization/flatten/fully connected..




## Convolutional Neural Networks (CNN)

In a machine learning model we can distinguish:

- ▶ **Parameters** – learned during training (e.g. weights of the connections between neurons)
- ▶ **Hyperparameters** – fixed *a-priori* – define the structure of the network:

For a CNN:

1. number and type of layers
2. for each convolutional layer: filter size, number of channels/filters, stride (stride=2 in the figure below)
3. Additional layers: dropout/batch normalization/flatten/fully connected..




# Convolutional Neural Networks (CNN)

In a machine learning model we can distinguish:

- ▶ **Parameters** – learned during training (e.g. weights of the connections between neurons)
- ▶ **Hyperparameters** – fixed *a-priori* – define the structure of the network:

For a CNN:

1. number and type of layers
2. for each convolutional layer: filter size, number of channels/filters, stride (stride=2 in the figure below)
3. Additional layers: dropout/batch normalization/flatten/fully connected..





# Going deep

Key advances allowing for training deep networks:

1. **ReLU activation:** (Rectified Linear Unit) i.e. nonlinearity for a neuron:

$$y_i = f(W_{ik}x_k + b_i) \quad \underbrace{f(x) = \max(0, x)}_{ReLU}$$

2. **Dropout:** during training we turn off say 50% randomly chosen neurons from the previous layer

<https://arxiv.org/abs/1207.0580>

turns out to be absolutely crucial for training deep networks with lots of parameters and prevent overfitting..

3. **Batch normalization:** A layer which normalizes inputs for each training batch

<https://arxiv.org/abs/1502.03167>

increases training speed/also acts as a regularizer

4. **GPU computations** – computations on (NVIDIA) graphics cards...

# Going deep

Key advances allowing for training deep networks:

1. **ReLU activation:** (Rectified Linear Unit) i.e. nonlinearity for a neuron:

$$y_i = f(W_{ik}x_k + b_i) \quad \underbrace{f(x) = \max(0, x)}_{\text{ReLU}}$$

2. **Dropout:** during training we turn off say 50% randomly chosen neurons from the previous layer

<https://arxiv.org/abs/1207.0580>

turns out to be absolutely crucial for training deep networks with lots of parameters and prevent overfitting..

3. **Batch normalization:** A layer which normalizes inputs for each training batch

<https://arxiv.org/abs/1502.03167>

increases training speed/also acts as a regularizer

4. **GPU computations** – computations on (NVIDIA) graphics cards...

# Going deep

Key advances allowing for training deep networks:

1. **ReLU activation:** (Rectified Linear Unit) i.e. nonlinearity for a neuron:

$$y_i = f(W_{ik}x_k + b_i) \quad \underbrace{f(x) = \max(0, x)}_{\text{ReLU}}$$

2. **Dropout:** during training we turn off say 50% randomly chosen neurons from the previous layer

<https://arxiv.org/abs/1207.0580>

turns out to be absolutely crucial for training deep networks with lots of parameters and prevent overfitting..

3. **Batch normalization:** A layer which normalizes inputs for each training batch

<https://arxiv.org/abs/1502.03167>

increases training speed/also acts as a regularizer

4. **GPU computations** – computations on (NVIDIA) graphics cards...

# Going deep

Key advances allowing for training deep networks:

1. **ReLU activation:** (Rectified Linear Unit) i.e. nonlinearity for a neuron:

$$y_i = f(W_{ik}x_k + b_i) \quad \underbrace{f(x) = \max(0, x)}_{\text{ReLU}}$$

2. **Dropout:** during training we turn off say 50% randomly chosen neurons from the previous layer

<https://arxiv.org/abs/1207.0580>

turns out to be absolutely crucial for training deep networks with lots of parameters and prevent overfitting..

3. **Batch normalization:** A layer which normalizes inputs for each training batch

<https://arxiv.org/abs/1502.03167>

increases training speed/also acts as a regularizer

4. **GPU computations** – computations on (NVIDIA) graphics cards...

# Going deep

Key advances allowing for training deep networks:

1. **ReLU activation:** (Rectified Linear Unit) i.e. nonlinearity for a neuron:

$$y_i = f(W_{ik}x_k + b_i) \quad \underbrace{f(x) = \max(0, x)}_{\text{ReLU}}$$

2. **Dropout:** during training we turn off say 50% randomly chosen neurons from the previous layer

<https://arxiv.org/abs/1207.0580>

turns out to be absolutely crucial for training deep networks with lots of parameters and prevent overfitting..

3. **Batch normalization:** A layer which normalizes inputs for each training batch

<https://arxiv.org/abs/1502.03167>

increases training speed/also acts as a regularizer

4. **GPU computations** – computations on (NVIDIA) graphics cards...

## Going deep

Key advances allowing for training deep networks:

1. **ReLU activation:** (Rectified Linear Unit) i.e. nonlinearity for a neuron:

$$y_i = f(W_{ik}x_k + b_i) \quad \underbrace{f(x) = \max(0, x)}_{\text{ReLU}}$$

2. **Dropout:** during training we turn off say 50% randomly chosen neurons from the previous layer

<https://arxiv.org/abs/1207.0580>

turns out to be absolutely crucial for training deep networks with lots of parameters and prevent overfitting..

3. **Batch normalization:** A layer which normalizes inputs for each training batch

<https://arxiv.org/abs/1502.03167>

increases training speed/also acts as a regularizer

4. **GPU computations** – computations on (NVIDIA) graphics cards...

## Going deep

Key advances allowing for training deep networks:

1. **ReLU activation:** (Rectified Linear Unit) i.e. nonlinearity for a neuron:

$$y_i = f(W_{ik}x_k + b_i) \quad \underbrace{f(x) = \max(0, x)}_{\text{ReLU}}$$

2. **Dropout:** during training we turn off say 50% randomly chosen neurons from the previous layer

<https://arxiv.org/abs/1207.0580>

turns out to be absolutely crucial for training deep networks with lots of parameters and prevent overfitting..

3. **Batch normalization:** A layer which normalizes inputs for each training batch

<https://arxiv.org/abs/1502.03167>

increases training speed/also acts as a regularizer

4. **GPU computations** – computations on (NVIDIA) graphics cards...

## CNN example: MNIST

Example from the first journal club using keras:

```
model = Sequential()
model.add(Convolution2D(32, 3, 3,
                        border_mode='valid',
                        input_shape=(1, 28, 28)))
model.add(Activation('relu'))
model.add(Convolution2D(32, 3, 3))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(128))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(10))
model.add(Activation('softmax'))
```

---

reaches 99.0% accuracy (MNIST subset)

CPU computation time: 14min

GPU computation time: 42s



## CNN example: MNIST

Example from the first journal club using keras:

```
model = Sequential()
model.add(Convolution2D(32, 3, 3,
                        border_mode='valid',
                        input_shape=(1, 28, 28)))
model.add(Activation('relu'))
model.add(Convolution2D(32, 3, 3))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(128))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(10))
model.add(Activation('softmax'))
```

---

reaches 99.0% accuracy (MNIST subset)

CPU computation time: 14min

GPU computation time: 42s

## CNN example: MNIST

Example from the first journal club using keras:

```
model = Sequential()
model.add(Convolution2D(32, 3, 3,
                        border_mode='valid',
                        input_shape=(1, 28, 28)))
model.add(Activation('relu'))
model.add(Convolution2D(32, 3, 3))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(128))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(10))
model.add(Activation('softmax'))
```

---

reaches 99.0% accuracy (MNIST subset)

CPU computation time: **14min**

GPU computation time: 42s

## CNN example: MNIST

Example from the first journal club using keras:

```
model = Sequential()
model.add(Convolution2D(32, 3, 3,
                        border_mode='valid',
                        input_shape=(1, 28, 28)))
model.add(Activation('relu'))
model.add(Convolution2D(32, 3, 3))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(128))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(10))
model.add(Activation('softmax'))
```

---

reaches 99.0% accuracy (MNIST subset)

CPU computation time: 14min

GPU computation time: 42s

## Image classification datasets – benchmarks for new CNN architectures

- ▶ **MNIST**: The *Hello World* of Machine Learning...  
60000+10000 28x28 hand written digits
- ▶ **CIFAR-10** and **CIFAR-100**  
60000 32x32 colour images in 10 classes (resp. 100 classes)

## Image classification datasets – benchmarks for new CNN architectures

- ▶ **MNIST**: The *Hello World* of Machine Learning...  
60000+10000 28x28 hand written digits

8 6 2 0 2 3 6 9 9 7

- ▶ **CIFAR-10** and **CIFAR-100**  
60000 32x32 colour images in 10 classes (resp. 100 classes)

## Image classification datasets – benchmarks for new CNN architectures

- ▶ **MNIST**: The *Hello World* of Machine Learning...  
60000+10000 28x28 hand written digits

8 6 2 0 2 3 6 9 9 7

- ▶ **CIFAR-10** and **CIFAR-100**  
60000 32x32 colour images in 10 classes (resp. 100 classes)

## Image classification datasets – benchmarks for new CNN architectures

- ▶ **MNIST**: The *Hello World* of Machine Learning...  
60000+10000 28x28 hand written digits

8 6 2 0 2 3 6 9 9 7

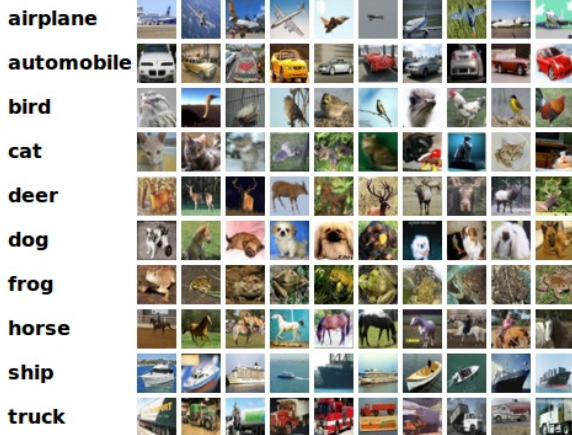
- ▶ **CIFAR-10** and **CIFAR-100**  
60000 32x32 colour images in 10 classes (resp. 100 classes)

## Image classification datasets – benchmarks for new CNN architectures

- **MNIST**: The *Hello World* of Machine Learning...  
60000+10000  $28 \times 28$  hand written digits

8 6 2 0 2 3 6 9 9 7

- **CIFAR-10** and **CIFAR-100**  
60000  $32 \times 32$  colour images in 10 classes (resp. 100 classes)





## Image classification datasets – benchmarks for new CNN architectures

- ▶ **ImageNet**: state of the art...

1.2 million images (256x256+) in 1000 classes

t-SNE visualization

<http://stanford.io/2iEFn9L>

## Image classification datasets – benchmarks for new CNN architectures

- ▶ **ImageNet**: state of the art...

1.2 million images (256x256+) in 1000 classes

t-SNE visualization

<http://stanford.io/2iEFn9L>

## Image classification datasets – benchmarks for new CNN architectures

- ▶ **ImageNet**: state of the art...

1.2 million images (256x256+) in 1000 classes

t-SNE visualization

<http://stanford.io/2iEFn9L>

# Image classification datasets – benchmarks for new CNN architectures

- **ImageNet**: state of the art...

1.2 million images (256x256+) in 1000 classes

t-SNE visualization

<http://stanford.io/2iEFn9L>



## Data augmentation

- ▶ In order to train deep neural networks one has to have lots of data
- ▶ This is necessary to
  1. avoid overfitting
  2. tell the network what are the inessential features (which are intuitive for us)
- ▶ For this reason one employs **data augmentation**  
— construct additional training examples by flipping/rotating/  
choosing different crops etc.

## Data augmentation

- ▶ In order to train deep neural networks one has to have lots of data
- ▶ This is necessary to
  1. avoid overfitting
  2. tell the network what are the inessential features (which are intuitive for us)
- ▶ For this reason one employs **data augmentation**  
— construct additional training examples by flipping/rotating/  
choosing different crops etc.

## Data augmentation

- ▶ In order to train deep neural networks one has to have lots of data
- ▶ This is necessary to
  1. avoid overfitting
  2. tell the network what are the inessential features (which are intuitive for us)
- ▶ For this reason one employs **data augmentation**  
— construct additional training examples by flipping/rotating/  
choosing different crops etc.

## Data augmentation

- ▶ In order to train deep neural networks one has to have lots of data
- ▶ This is necessary to
  1. avoid overfitting
  2. tell the network what are the inessential features (which are intuitive for us)
- ▶ For this reason one employs **data augmentation**  
— construct additional training examples by flipping/rotating/  
choosing different crops etc.



## Data augmentation

- ▶ In order to train deep neural networks one has to have lots of data
- ▶ This is necessary to
  1. avoid overfitting
  2. tell the network what are the inessential features (which are intuitive for us)
- ▶ For this reason one employs **data augmentation**
  - construct additional training examples by flipping/rotating/choosing different crops etc.

## Data augmentation

- ▶ In order to train deep neural networks one has to have lots of data
- ▶ This is necessary to
  1. avoid overfitting
  2. tell the network what are the inessential features (which are intuitive for us)
- ▶ For this reason one employs **data augmentation**  
— construct additional training examples by flipping/rotating/  
choosing different crops etc.


## Image classification datasets – benchmarks for new CNN architectures

### Some key results on ImageNet:

- ▶ AlexNet result in 2012 was a breakthrough — 2<sup>nd</sup> result: 26.2%
- ▶ The progress was **not** only due to brute force!  
AlexNet ~ 60M parameters, VGG ~ 180M but GoogLeNet only 5M  
residual network also has a new architectural ingredient
- ▶ Human performance: 5.1% error (see <http://bit.ly/2f45u9i>)
- ▶ Top-1 errors are much larger (17% – 20%+)
- ▶ 2016: GoogLeNet Inception V4 (3.08%)

# Image classification datasets – benchmarks for new CNN architectures


## Some key results on ImageNet:



- ▶ AlexNet result in 2012 was a breakthrough — 2<sup>nd</sup> result: 26.2%
- ▶ The progress was **not** only due to brute force!  
AlexNet ~ 60M parameters, VGG ~ 180M but GoogLeNet only 5M  
residual network also has a new architectural ingredient
- ▶ Human performance: 5.1% error (see <http://bit.ly/2f45u9i>)
- ▶ Top-1 errors are much larger (17% – 20%+)
- ▶ 2016: GoogLeNet Inception V4 (3.08%)

## Image classification datasets – benchmarks for new CNN architectures


### Some key results on ImageNet:



- ▶ AlexNet result in 2012 was a breakthrough — 2<sup>nd</sup> result: 26.2%
- ▶ The progress was **not** only due to brute force!  
AlexNet ~ 60M parameters, VGG ~ 180M but GoogLeNet only 5M  
residual network also has a new architectural ingredient
- ▶ Human performance: 5.1% error (see <http://bit.ly/2f45u9i>)
- ▶ Top-1 errors are much larger (17% – 20%+)
- ▶ 2016: GoogLeNet Inception V4 (3.08%)

# Image classification datasets – benchmarks for new CNN architectures


## Some key results on ImageNet:



- ▶ AlexNet result in 2012 was a breakthrough —  $2^{nd}$  result: 26.2%
- ▶ The progress was **not** only due to brute force!  
AlexNet  $\sim 60M$  parameters, VGG  $\sim 180M$  but GoogLeNet only 5M  
residual network also has a new architectural ingredient
- ▶ Human performance: 5.1% error (see <http://bit.ly/2f45u9i>)
- ▶ Top-1 errors are much larger (17% – 20%+)
- ▶ 2016: GoogLeNet Inception V4 (3.08%)

# Image classification datasets – benchmarks for new CNN architectures


## Some key results on ImageNet:



- ▶ AlexNet result in 2012 was a breakthrough —  $2^{nd}$  result: 26.2%
- ▶ The progress was **not** only due to brute force!  
AlexNet  $\sim 60M$  parameters, VGG  $\sim 180M$  but GoogLeNet only 5M  
residual network also has a new architectural ingredient
- ▶ Human performance: 5.1% error (see <http://bit.ly/2f45u9i>)
- ▶ Top-1 errors are much larger (17% – 20%+)
- ▶ 2016: GoogLeNet Inception V4 (3.08%)

## Image classification datasets – benchmarks for new CNN architectures

### Some key results on ImageNet:




- ▶ AlexNet result in 2012 was a breakthrough —  $2^{nd}$  result: 26.2%
- ▶ The progress was **not** only due to brute force!  
AlexNet  $\sim 60M$  parameters, VGG  $\sim 180M$  but GoogLeNet only  $5M$   
residual network also has a new architectural ingredient
- ▶ Human performance: 5.1% error (see <http://bit.ly/2f45u9i>)
- ▶ Top-1 errors are much larger (17% – 20%+)
- ▶ 2016: GoogLeNet Inception V4 (3.08%)



# Image classification datasets – benchmarks for new CNN architectures


## Some key results on ImageNet:



- ▶ AlexNet result in 2012 was a breakthrough —  $2^{nd}$  result: 26.2%
- ▶ The progress was **not** only due to brute force!  
AlexNet  $\sim 60M$  parameters, VGG  $\sim 180M$  but GoogLeNet only  $5M$   
residual network also has a new architectural ingredient
- ▶ Human performance: 5.1% error (see <http://bit.ly/2f45u9i>)
- ▶ Top-1 errors are much larger (17% – 20%+)
- ▶ 2016: GoogLeNet Inception V4 (3.08%)

# Image classification datasets – benchmarks for new CNN architectures


## Some key results on ImageNet:



- ▶ AlexNet result in 2012 was a breakthrough —  $2^{nd}$  result: 26.2%
- ▶ The progress was **not** only due to brute force!  
AlexNet  $\sim 60M$  parameters, VGG  $\sim 180M$  but GoogLeNet only  $5M$   
residual network also has a new architectural ingredient
- ▶ Human performance: 5.1% error (see <http://bit.ly/2f45u9i>)
- ▶ Top-1 errors are much larger (17% – 20%+)
- ▶ 2016: GoogLeNet Inception V4 (3.08%)

# Image classification datasets – benchmarks for new CNN architectures

## Some key results on ImageNet:



- ▶ AlexNet result in 2012 was a breakthrough —  $2^{nd}$  result: 26.2%
- ▶ The progress was **not** only due to brute force!  
AlexNet  $\sim 60M$  parameters, VGG  $\sim 180M$  but GoogLeNet only  $5M$   
residual network also has a new architectural ingredient
- ▶ Human performance: 5.1% error (see <http://bit.ly/2f45u9i>)
- ▶ Top-1 errors are much larger (17% – 20%+)
- ▶ 2016: GoogLeNet Inception V4 (3.08%)

## Recent CNN architectures

<https://arxiv.org/abs/1409.4842>

<https://arxiv.org/abs/1409.4842>

- ▶ Classical CNN's are a sequence of convolution and pooling layers
- ▶ Two ideas:
  1. Allow filters of different sizes and concatenate
  2. Each convolution takes as input  $n_{channels} \times N_x \times N_y$ .  
First perform pointwise convolution over channels ( $1 \times 1$  filter) and then spatial convolution
- ▶ This significantly reduces the number of parameters..
- ▶ GoogLeNet stacks multiple Inception modules: 22 layer deep

<https://arxiv.org/abs/1409.4842>

- ▶ Classical CNN's are a sequence of convolution and pooling layers
- ▶ Two ideas:
  1. Allow filters of different sizes and concatenate
  2. Each convolution takes as input  $n_{channels} \times N_x \times N_y$ .  
First perform pointwise convolution over channels ( $1 \times 1$  filter) and then spatial convolution
- ▶ This significantly reduces the number of parameters..
- ▶ GoogLeNet stacks multiple Inception modules: 22 layer deep

<https://arxiv.org/abs/1409.4842>



<https://arxiv.org/abs/1409.4842>


- ▶ Classical CNN's are a sequence of convolution and pooling layers
  - ▶ Two ideas:
    1. Allow filters of different sizes and concatenate
    2. Each convolution takes as input  $n_{channels} \times N_x \times N_y$ .  
First perform pointwise convolution over channels ( $1 \times 1$  filter) and then spatial convolution
- 
- ▶ This significantly reduces the number of parameters..
  - ▶ GoogLeNet stacks multiple Inception modules: 22 layer deep

<https://arxiv.org/abs/1409.4842>

- ▶ Classical CNN's are a sequence of convolution and pooling layers
  - ▶ Two ideas:
    1. Allow filters of different sizes and concatenate
    2. Each convolution takes as input  $n_{channels} \times N_x \times N_y$ .  
First perform pointwise convolution over channels ( $1 \times 1$  filter) and then spatial convolution
- 
- ▶ This significantly reduces the number of parameters..
  - ▶ GoogLeNet stacks multiple Inception modules: 22 layer deep

<https://arxiv.org/abs/1409.4842>


- ▶ Classical CNN's are a sequence of convolution and pooling layers
- ▶ Two ideas:
  1. Allow filters of different sizes and concatenate
  2. Each convolution takes as input  $n_{channels} \times N_x \times N_y$ .  
First perform pointwise convolution over channels ( $1 \times 1$  filter) and then spatial convolution



- ▶ This significantly reduces the number of parameters..
- ▶ GoogLeNet stacks multiple Inception modules: 22 layer deep

<https://arxiv.org/abs/1409.4842>

- ▶ Classical CNN's are a sequence of convolution and pooling layers
- ▶ Two ideas:
  1. Allow filters of different sizes and concatenate
  2. Each convolution takes as input  $n_{channels} \times N_x \times N_y$ .  
First perform pointwise convolution over channels ( $1 \times 1$  filter) and then spatial convolution



- ▶ This significantly reduces the number of parameters..
- ▶ GoogLeNet stacks multiple Inception modules: 22 layer deep

F. Chollet <https://arxiv.org/abs/1610.02357>

- ▶ One can take factorization to an extreme...
- ▶ Factor all convolutions into a pointwise and spatial convolution
- ▶ This implemented in TensorFlow/Keras as first performing spatial convolution for each channel separately and then projecting pointwise by a  $1 \times 1$  convolution
- ▶ Works better than Inception V3 with same number of parameters
- ▶ Use also an internal Google dataset: JFT: 350 million high resolution images annotated with 17000 classes (mutli-label)
- ▶ Models trained on 60 NVIDIA GPU's

F. Chollet <https://arxiv.org/abs/1610.02357>

- ▶ One can take factorization to an extreme...
- ▶ Factor all convolutions into a pointwise and spatial convolution
- ▶ This implemented in TensorFlow/Keras as first performing spatial convolution for each channel separately and then projecting pointwise by a  $1 \times 1$  convolution
- ▶ Works better than Inception V3 with same number of parameters
- ▶ Use also an internal Google dataset: JFT: 350 million high resolution images annotated with 17000 classes (mutli-label)
- ▶ Models trained on 60 NVIDIA GPU's

F. Chollet <https://arxiv.org/abs/1610.02357>

- ▶ One can take factorization to an extreme...
- ▶ Factor all convolutions into a pointwise and spatial convolution
- ▶ This implemented in TensorFlow/Keras as first performing spatial convolution for each channel separately and then projecting pointwise by a  $1 \times 1$  convolution
- ▶ Works better than Inception V3 with same number of parameters
- ▶ Use also an internal Google dataset: JFT: 350 million high resolution images annotated with 17000 classes (mutli-label)
- ▶ Models trained on 60 NVIDIA GPU's

F. Chollet <https://arxiv.org/abs/1610.02357>

- ▶ One can take factorization to an extreme...
- ▶ Factor all convolutions into a pointwise and spatial convolution
- ▶ This implemented in TensorFlow/Keras as first performing spatial convolution for each channel separately and then projecting pointwise by a  $1 \times 1$  convolution
- ▶ Works better than Inception V3 with same number of parameters
- ▶ Use also an internal Google dataset: JFT: 350 million high resolution images annotated with 17000 classes (mutli-label)
- ▶ Models trained on 60 NVIDIA GPU's



F. Chollet <https://arxiv.org/abs/1610.02357>

- ▶ One can take factorization to an extreme...
- ▶ Factor all convolutions into a pointwise and spatial convolution
- ▶ This implemented in TensorFlow/Keras as first performing spatial convolution for each channel separately and then projecting pointwise by a  $1 \times 1$  convolution
- ▶ Works better than Inception V3 with same number of parameters
- ▶ Use also an internal Google dataset: JFT: 350 million high resolution images annotated with 17000 classes (mutli-label)
- ▶ Models trained on 60 NVIDIA GPU's

F. Chollet <https://arxiv.org/abs/1610.02357>

- ▶ One can take factorization to an extreme...
- ▶ Factor all convolutions into a pointwise and spatial convolution
- ▶ This implemented in TensorFlow/Keras as first performing spatial convolution for each channel separately and then projecting pointwise by a  $1 \times 1$  convolution
- ▶ Works better than Inception V3 with same number of parameters
- ▶ Use also an internal Google dataset: JFT: 350 million high resolution images annotated with 17000 classes (mutli-label)
- ▶ Models trained on 60 NVIDIA GPU's

F. Chollet <https://arxiv.org/abs/1610.02357>

- ▶ One can take factorization to an extreme...
- ▶ Factor all convolutions into a pointwise and spatial convolution
- ▶ This implemented in TensorFlow/Keras as first performing spatial convolution for each channel separately and then projecting pointwise by a  $1 \times 1$  convolution
- ▶ Works better than Inception V3 with same number of parameters
- ▶ Use also an internal Google dataset: JFT: 350 million high resolution images annotated with 17000 classes (mutli-label)
- ▶ Models trained on **60** NVIDIA GPU's

<https://arxiv.org/abs/1512.03385>

<https://arxiv.org/abs/1603.05027>

- ▶ In what way are very deep networks difficult to train?  
**Not** due to overfitting! – have difficulty learning the training set
- ▶ But in principle one could find a trivial 56-layer network which is as good as a 20 layer one — just take a 20 layer one and add 36 layers which realize an identity mapping..
- ▶ **Conclusion:** in conventional CNN's it is difficult to make deep layers learn the identity mapping!

<https://arxiv.org/abs/1512.03385>

<https://arxiv.org/abs/1603.05027>

- ▶ In what way are *very* deep networks difficult to train?  
**Not** due to overfitting! – have difficulty learning the training set
- ▶ But in principle one could find a trivial 56-layer network which is as good as a 20 layer one — just take a 20 layer one and add 36 layers which realize an identity mapping..
- ▶ **Conclusion:** in conventional CNN's it is difficult to make deep layers learn the identity mapping!

<https://arxiv.org/abs/1512.03385>

<https://arxiv.org/abs/1603.05027>

- ▶ In what way are *very* deep networks difficult to train?  
**Not** due to overfitting! – have difficulty learning the training set
- ▶ But in principle one could find a trivial 56-layer network which is as good as a 20 layer one — just take a 20 layer one and add 36 layers which realize an identity mapping..
- ▶ **Conclusion:** in conventional CNN's it is difficult to make deep layers learn the identity mapping!

<https://arxiv.org/abs/1512.03385>


<https://arxiv.org/abs/1603.05027>

- ▶ In what way are *very* deep networks difficult to train?  
**Not** due to overfitting! – have difficulty learning the training set
- ▶ But in principle one could find a trivial 56-layer network which is as good as a 20 layer one — just take a 20 layer one and add 36 layers which realize an identity mapping..
- ▶ **Conclusion:** in conventional CNN's it is difficult to make deep layers learn the identity mapping!

<https://arxiv.org/abs/1512.03385>

<https://arxiv.org/abs/1603.05027>

- In what way are *very* deep networks difficult to train?  
**Not** due to overfitting! – have difficulty learning the training set




- But in principle one could find a trivial 56-layer network which is as good as a 20 layer one — just take a 20 layer one and add 36 layers which realize an identity mapping..
- **Conclusion:** in conventional CNN's it is difficult to make deep layers learn the identity mapping!



<https://arxiv.org/abs/1512.03385>

<https://arxiv.org/abs/1603.05027>

- In what way are *very* deep networks difficult to train?  
**Not** due to overfitting! – have difficulty learning the training set




- But in principle one could find a trivial 56-layer network which is as good as a 20 layer one — just take a 20 layer one and add 36 layers which realize an identity mapping..
- **Conclusion:** in conventional CNN's it is difficult to make deep layers learn the identity mapping!

<https://arxiv.org/abs/1512.03385>

<https://arxiv.org/abs/1603.05027>

- ▶ In what way are *very* deep networks difficult to train?  
**Not** due to overfitting! – have difficulty learning the training set




- ▶ But in principle one could find a trivial 56-layer network which is as good as a 20 layer one — just take a 20 layer one and add 36 layers which realize an identity mapping..
- ▶ **Conclusion:** in conventional CNN's it is difficult to make deep layers learn the identity mapping!

<https://arxiv.org/abs/1512.03385>

<https://arxiv.org/abs/1603.05027>

- ▶ In what way are *very* deep networks difficult to train?  
**Not** due to overfitting! – have difficulty learning the training set



- ▶ But in principle one could find a trivial 56-layer network which is as good as a 20 layer one — just take a 20 layer one and add 36 layers which realize an identity mapping..
- ▶ **Conclusion:** in conventional CNN's it is difficult to make deep layers learn the identity mapping!


**Proposal:** Learn deviations from identity rather than the full mapping...

- ▶ Stacking such residual units gives CNN's with many layers: 200 on ImageNet and even 1001 on CIFAR!
- ▶ Various variations of this idea have been tested in the second paper (this is the latest version)
- ▶ Residual networks have become the state-of-the-art
- ▶ The residual network ideas are being incorporated into other architectures (like Inception, Fully Convolutional Networks (FCN) for segmentation etc.)

**Proposal:** Learn deviations from identity rather than the full mapping...


- ▶ Stacking such residual units gives CNN's with many layers: 200 on ImageNet and even 1001 on CIFAR!
- ▶ Various variations of this idea have been tested in the second paper (this is the latest version)
- ▶ Residual networks have become the state-of-the-art
- ▶ The residual network ideas are being incorporated into other architectures (like Inception, Fully Convolutional Networks (FCN) for segmentation etc.)

**Proposal:** Learn deviations from identity rather than the full mapping...




- ▶ Stacking such residual units gives CNN's with many layers: 200 on ImageNet and even 1001 on CIFAR!
- ▶ Various variations of this idea have been tested in the second paper (this is the latest version)
- ▶ Residual networks have become the state-of-the-art
- ▶ The residual network ideas are being incorporated into other architectures (like Inception, Fully Convolutional Networks (FCN) for segmentation etc.)

**Proposal:** Learn deviations from identity rather than the full mapping...



- ▶ Stacking such residual units gives CNN's with many layers: 200 on ImageNet and even 1001 on CIFAR!
- ▶ Various variations of this idea have been tested in the second paper (this is the latest version)
- ▶ Residual networks have become the state-of-the-art
- ▶ The residual network ideas are being incorporated into other architectures (like Inception, Fully Convolutional Networks (FCN) for segmentation etc.)


**Proposal:** Learn deviations from identity rather than the full mapping...



- ▶ Stacking such residual units gives CNN's with many layers: 200 on ImageNet and even 1001 on CIFAR!
- ▶ Various variations of this idea have been tested in the second paper (this is the latest version)
- ▶ Residual networks have become the state-of-the-art
- ▶ The residual network ideas are being incorporated into other architectures (like Inception, Fully Convolutional Networks (FCN) for segmentation etc.)




**Proposal:** Learn deviations from identity rather than the full mapping...



- ▶ Stacking such residual units gives CNN's with many layers: 200 on ImageNet and even 1001 on CIFAR!
- ▶ Various variations of this idea have been tested in the second paper (this is the latest version)
- ▶ Residual networks have become the state-of-the-art
- ▶ The residual network ideas are being incorporated into other architectures (like Inception, Fully Convolutional Networks (FCN) for segmentation etc.)

**Proposal:** Learn deviations from identity rather than the full mapping...



- ▶ Stacking such residual units gives CNN's with many layers: 200 on ImageNet and even 1001 on CIFAR!
- ▶ Various variations of this idea have been tested in the second paper (this is the latest version)
- ▶ Residual networks have become the state-of-the-art
- ▶ The residual network ideas are being incorporated into other architectures (like Inception, Fully Convolutional Networks (FCN) for segmentation etc.)

<https://arxiv.org/abs/1608.06993>

- ▶ Residual networks (and earlier more complicated *Highway networks* + other ideas: *stochastic depth*, *FractalNets*) provide multiple paths for data to flow through the networks
  - ▶ Densely connected convolutional networks (DensNet) exploit this idea
- 
- ▶ The outputs of the layers in each block are concatenated (and not summed)
  - ▶ In the transition layers the authors reduce the number of features while in the convolution layers introduce  $1 \times 1$  convolutions (bottleneck layers)

<https://arxiv.org/abs/1608.06993>


- ▶ Residual networks (and earlier more complicated *Highway networks* + other ideas: *stochastic depth*, *FractalNets*) provide multiple paths for data to flow through the networks
- ▶ Densely connected convolutional networks (DensNet) exploit this idea
- ▶ The outputs of the layers in each block are concatenated (and not summed)
- ▶ In the transition layers the authors reduce the number of features while in the convolution layers introduce  $1 \times 1$  convolutions (bottleneck layers)

<https://arxiv.org/abs/1608.06993>

- ▶ Residual networks (and earlier more complicated *Highway networks* + other ideas: *stochastic depth*, *FractalNets*) provide multiple paths for data to flow through the networks
  - ▶ Densely connected convolutional networks (DensNet) exploit this idea
- 
- ▶ The outputs of the layers in each block are concatenated (and not summed)
  - ▶ In the transition layers the authors reduce the number of features while in the convolution layers introduce  $1 \times 1$  convolutions (bottleneck layers)

<https://arxiv.org/abs/1608.06993>


- ▶ Residual networks (and earlier more complicated *Highway networks* + other ideas: *stochastic depth*, *FractalNets*) provide multiple paths for data to flow through the networks
- ▶ Densely connected convolutional networks (DensNet) exploit this idea



- ▶ The outputs of the layers in each block are concatenated (and not summed)
- ▶ In the transition layers the authors reduce the number of features while in the convolution layers introduce  $1 \times 1$  convolutions (bottleneck layers)

<https://arxiv.org/abs/1608.06993>


- ▶ Residual networks (and earlier more complicated *Highway networks* + other ideas: *stochastic depth*, *FractalNets*) provide multiple paths for data to flow through the networks
- ▶ Densely connected convolutional networks (DensNet) exploit this idea



- ▶ The outputs of the layers in each block are concatenated (and not summed)
- ▶ In the transition layers the authors reduce the number of features while in the convolution layers introduce  $1 \times 1$  convolutions (bottleneck layers)

<https://arxiv.org/abs/1608.06993>

- ▶ Residual networks (and earlier more complicated *Highway networks* + other ideas: *stochastic depth*, *FractalNets*) provide multiple paths for data to flow through the networks
- ▶ Densely connected convolutional networks (DensNet) exploit this idea




- ▶ The outputs of the layers in each block are concatenated (and not summed)
- ▶ In the transition layers the authors reduce the number of features while in the convolution layers introduce  $1 \times 1$  convolutions (bottleneck layers)




<https://arxiv.org/abs/1608.06993>

- ▶ Residual networks (and earlier more complicated *Highway networks* + other ideas: *stochastic depth*, *FractalNets*) provide multiple paths for data to flow through the networks
- ▶ Densely connected convolutional networks (DensNet) exploit this idea



- ▶ The outputs of the layers in each block are concatenated (and not summed)
- ▶ In the transition layers the authors reduce the number of features while in the convolution layers introduce  $1 \times 1$  convolutions (bottleneck layers)



Top-1 error on ImageNet


## Interlude

## Intriguing properties of neural networks

<https://arxiv.org/abs/1312.6199>

- ▶ On the left: images which correctly classified by AlexNet (deep CNN trained on ImageNet)
- ▶ Center: perturbations magnified  $10\times$
- ▶ On the right: images classified by the same net as **ostrich!!!**


# Intriguing properties of neural networks



<https://arxiv.org/abs/1312.6199>

- ▶ On the left: images which correctly classified by AlexNet (deep CNN trained on ImageNet)
- ▶ Center: perturbations magnified  $10\times$
- ▶ On the right: images classified by the same net as **ostrich!!!**


# Intriguing properties of neural networks



<https://arxiv.org/abs/1312.6199>

- ▶ On the left: images which correctly classified by AlexNet (deep CNN trained on ImageNet)
- ▶ Center: perturbations magnified 10 $\times$
- ▶ On the right: images classified by the same net as **ostrich!!!**


# Intriguing properties of neural networks



<https://arxiv.org/abs/1312.6199>

- ▶ On the left: images which correctly classified by AlexNet (deep CNN trained on ImageNet)
- ▶ Center: perturbations magnified  $10\times$
- ▶ On the right: images classified by the same net as **ostrich!!!**

# Intriguing properties of neural networks



<https://arxiv.org/abs/1312.6199>

- ▶ On the left: images which correctly classified by AlexNet (deep CNN trained on ImageNet)
- ▶ Center: perturbations magnified  $10\times$
- ▶ On the right: images classified by the same net as **ostrich!!!**

# Deep Neural Networks are Easily Fooled: High Confidence Predictions for Unrecognizable Images

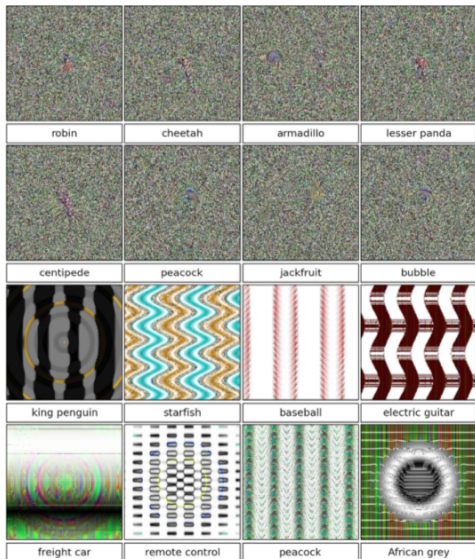
<https://arxiv.org/abs/1412.1897>

- ▶ Images which are classified by AlexNet (deep CNN trained on ImageNet) with more than 99.6% confidence
- ▶ These examples were constructed by a kind of evolutionary procedure...

*Our results shed light on interesting differences between human vision and current DNNs, and raise questions about the generality of DNN computer vision.*



# Deep Neural Networks are Easily Fooled: High Confidence Predictions for Unrecognizable Images



<https://arxiv.org/abs/1412.1897>

- Images which are classified by AlexNet (deep CNN trained on ImageNet) with more than 99.6% confidence
- These examples were constructed by a kind of evolutionary procedure...


*Our results shed light on interesting differences between human vision and current DNNs, and raise questions about the generality of DNN computer vision.*

# Deep Neural Networks are Easily Fooled: High Confidence Predictions for Unrecognizable Images

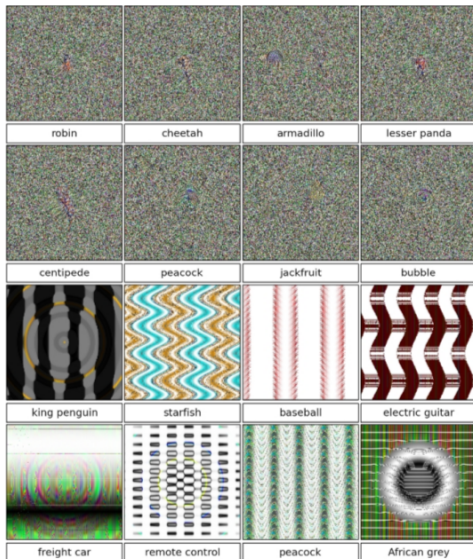
<https://arxiv.org/abs/1412.1897>

- ▶ Images which are classified by AlexNet (deep CNN trained on ImageNet) with more than 99.6% confidence
- ▶ These examples were constructed by a kind of evolutionary procedure...

*Our results shed light on interesting differences between human vision and current DNNs, and raise questions about the generality of DNN computer vision.*



# Deep Neural Networks are Easily Fooled: High Confidence Predictions for Unrecognizable Images



<https://arxiv.org/abs/1412.1897>

- ▶ Images which are classified by AlexNet (deep CNN trained on ImageNet) with more than 99.6% confidence
- ▶ These examples were constructed by a kind of evolutionary procedure...

*Our results shed light on interesting differences between human vision and current DNNs, and raise questions about the generality of DNN computer vision.*

Goodfellow, Shlens, Szegedy Explaining and Harnessing Adversarial Examples  
<https://arxiv.org/abs/1412.6572>

- ▶ Adversarial examples are due to *linear* nature (and high dimensionality) of the NN
- ▶ Not restricted to deep CNN's...

Goodfellow, Shlens, Szegedy Explaining and Harnessing Adversarial Examples  
<https://arxiv.org/abs/1412.6572>

- ▶ Adversarial examples are due to *linear* nature (and high dimensionality) of the NN
- ▶ Not restricted to deep CNN's...

Goodfellow, Shlens, Szegedy Explaining and Harnessing Adversarial Examples  
<https://arxiv.org/abs/1412.6572>

- ▶ Adversarial examples are due to *linear* nature (and high dimensionality) of the NN
- ▶ Not restricted to deep CNN's...

Goodfellow, Shlens, Szegedy Explaining and Harnessing Adversarial Examples  
<https://arxiv.org/abs/1412.6572>

- ▶ Adversarial examples are due to *linear* nature (and high dimensionality) of the NN
- ▶ Not restricted to deep CNN's...

# Image segmentation

## Image segmentation task:

Not only identify the class(es) of the objects in the image but locate them pixelwise

Very good set of lectures @Stanford '16 (slides and notes online!):

**CS321n: Convolutional Neural Networks for Visual Recognition**

<http://cs231n.stanford.edu/syllabus.html>



# Image segmentation

**Classification**




CAT

**Classification  
+ Localization**



CAT

**Object Detection**



CAT, DOG, DUCK

**Instance  
Segmentation**



CAT, DOG, DUCK

Single object

Multiple objects

## Image segmentation task:

Not only identify the class(es) of the objects in the image but locate them pixelwise

Very good set of lectures @Stanford '16 (slides and notes online!):  
CS321n: Convolutional Neural Networks for Visual Recognition

<http://cs231n.stanford.edu/syllabus.html>

# Image segmentation

**Classification**




CAT

**Classification  
+ Localization**



CAT

**Object Detection**



CAT, DOG, DUCK

**Instance  
Segmentation**



CAT, DOG, DUCK

Single object

Multiple objects

## Image segmentation task:

Not only identify the class(es) of the objects in the image but locate them pixelwise

Very good set of lectures @Stanford '16 (slides and notes online!):  
CS321n: Convolutional Neural Networks for Visual Recognition

<http://cs231n.stanford.edu/syllabus.html>

# Image segmentation

**Classification**




CAT

**Classification  
+ Localization**



CAT

**Object Detection**



CAT, DOG, DUCK

**Instance  
Segmentation**



CAT, DOG, DUCK

Single object

Multiple objects

## Image segmentation task:

Not only identify the class(es) of the objects in the image but locate them pixelwise

Very good set of lectures @Stanford '16 (slides and notes online!):

**CS321n: Convolutional Neural Networks for Visual Recognition**

<http://cs231n.stanford.edu/syllabus.html>


## Fully Convolutional Networks — U-Net

<https://arxiv.org/abs/1505.04597>

- ▶ Network used for biomedical image segmentation... (but generic segmentation architecture)
- ▶ Lots of papers, see e.g. <https://arxiv.org/abs/1611.09326>  
Y. Bengio et.al. *Fully convolutional DenseNets for semantic segmentation*

# Fully Convolutional Networks — U-Net

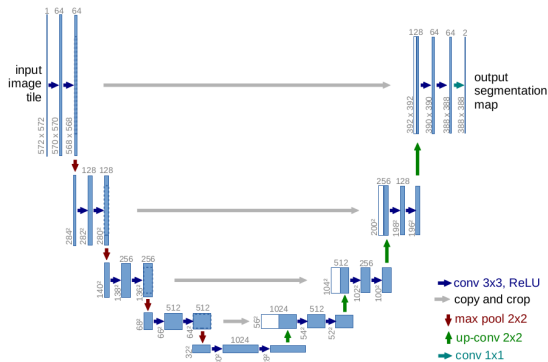
<https://arxiv.org/abs/1505.04597>



- ▶ Network used for biomedical image segmentation... (but generic segmentation architecture)
- ▶ Lots of papers, see e.g. <https://arxiv.org/abs/1611.09326>  
Y. Bengio et.al. *Fully convolutional DenseNets for semantic segmentation*

# Fully Convolutional Networks — U-Net


<https://arxiv.org/abs/1505.04597>



- ▶ Network used for biomedical image segmentation... (but generic segmentation architecture)
- ▶ Lots of papers, see e.g. <https://arxiv.org/abs/1611.09326>  
Y. Bengio et.al. *Fully convolutional DenseNets for semantic segmentation*

# Fully Convolutional Networks — U-Net

<https://arxiv.org/abs/1505.04597>

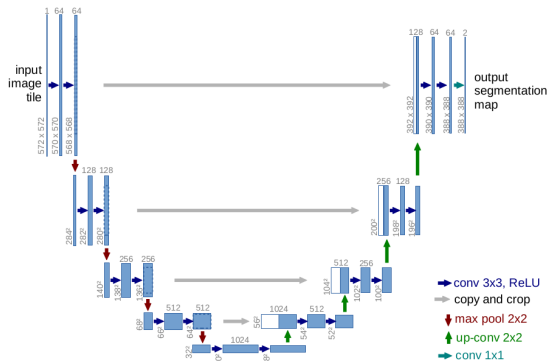


- ▶ Network used for biomedical image segmentation... (but generic segmentation architecture)

- ▶ Lots of papers, see e.g. <https://arxiv.org/abs/1611.09326>  
Y. Bengio et.al. *Fully convolutional DenseNets for semantic segmentation*

# Fully Convolutional Networks — U-Net

<https://arxiv.org/abs/1505.04597>



- ▶ Network used for biomedical image segmentation... (but generic segmentation architecture)
- ▶ Lots of papers, see e.g. <https://arxiv.org/abs/1611.09326>  
Y. Bengio et.al. *Fully convolutional DenseNets for semantic segmentation*



## Conclusions

- ▶ Convolutional networks are used primarily for images but can be also used for time-series data (1D convolutions) and video (3D convolution)
- ▶ Apart from classification, CNNs are also used e.g. for image segmentation
- ▶ GPU's are currently absolutely crucial for computation...
- ▶ Convolutional networks are still being developed and improved...  
... not only brute force but also novel architectures
- ▶ It seems that having multiple possibilities of traversing the network is very beneficial

## Conclusions

- ▶ Convolutional networks are used primarily for images but can be also used for time-series data (1D convolutions) and video (3D convolution)
- ▶ Apart from classification, CNNs are also used e.g. for image segmentation
- ▶ GPU's are currently absolutely crucial for computation...
- ▶ Convolutional networks are still being developed and improved...  
... not only brute force but also novel architectures
- ▶ It seems that having multiple possibilities of traversing the network is very beneficial

## Conclusions

- ▶ Convolutional networks are used primarily for images but can be also used for time-series data (1D convolutions) and video (3D convolution)
- ▶ Apart from classification, CNNs are also used e.g. for image segmentation
- ▶ GPU's are currently absolutely crucial for computation...
- ▶ Convolutional networks are still being developed and improved...  
... not only brute force but also novel architectures
- ▶ It seems that having multiple possibilities of traversing the network is very beneficial

## Conclusions

- ▶ Convolutional networks are used primarily for images but can be also used for time-series data (1D convolutions) and video (3D convolution)
- ▶ Apart from classification, CNNs are also used e.g. for image segmentation
- ▶ GPU's are currently absolutely crucial for computation...
- ▶ Convolutional networks are still being developed and improved...  
... not only brute force but also novel architectures
- ▶ It seems that having multiple possibilities of traversing the network is very beneficial

## Conclusions

- ▶ Convolutional networks are used primarily for images but can be also used for time-series data (1D convolutions) and video (3D convolution)
- ▶ Apart from classification, CNNs are also used e.g. for image segmentation
- ▶ GPU's are currently absolutely crucial for computation...
- ▶ Convolutional networks are still being developed and improved...  
... not only brute force but also novel architectures
- ▶ It seems that having multiple possibilities of traversing the network is very beneficial

## Conclusions

- ▶ Convolutional networks are used primarily for images but can be also used for time-series data (1D convolutions) and video (3D convolution)
- ▶ Apart from classification, CNNs are also used e.g. for image segmentation
- ▶ GPU's are currently absolutely crucial for computation...
- ▶ Convolutional networks are still being developed and improved...  
... not only brute force but also novel architectures
- ▶ It seems that having multiple possibilities of traversing the network is very beneficial

## Conclusions

- ▶ Convolutional networks are used primarily for images but can be also used for time-series data (1D convolutions) and video (3D convolution)
- ▶ Apart from classification, CNNs are also used e.g. for image segmentation
- ▶ GPU's are currently absolutely crucial for computation...
- ▶ Convolutional networks are still being developed and improved...  
... not only brute force but also novel architectures
- ▶ It seems that having multiple possibilities of traversing the network is very beneficial